

EPPS6323 Knowledge Mining

Spring 2026 Syllabus

Karl Ho, PhD.

January 27, 2026

1 Course Information

Schedule	Tuesday 4:00 PM - 6:45 PM
Location	FO 1.502
Semester	Spring 2026

1.1 Instructor Information

Instructor	Karl Ho, Ph.D.
Office	GR 3.203
Phone	972-883-2017
Email	kyho@utdallas.edu
Office Hours	Monday & Wednesday, 1:30 PM – 3:30 PM (by appointment)

1.2 Teaching Assistant

TA	Ariel Wang
Office	GR 3.318
Email	Ariel.Wang@UTDallas.edu
Office Hours	Wednesday & Thursday 1:30 PM – 3:30 PM (by appointment)

2 Course Overview

This course explores the interdisciplinary field of **knowledge mining**, integrating theoretical concepts from data mining, statistical learning, and machine learning, with a focus on leveraging AI, large language models (LLMs), and generative AI for knowledge discovery.

The 2026 edition of this course reflects the significant paradigm shift toward **agentic AI systems**—autonomous agents that can reason, plan, use tools, and execute multi-step tasks with minimal human intervention. Students will learn to work with AI as a **research partner**, developing skills in:

- Retrieval-Augmented Generation (RAG) architectures
- Multimodal knowledge extraction and knowledge graphs
- AI-assisted coding for scientific research
- Ethical AI deployment in research contexts

Students will gain hands-on experience with tools and techniques for analyzing and visualizing data, extracting insights, and applying ethical considerations in real-world research scenarios.

AI-Native Learning Approach

This course adopts an “AI as research partner” pedagogy. Students will progressively develop skills in human-AI collaboration, moving from AI-assisted work to AI-delegated research tasks with appropriate oversight.

3 Prerequisites

There is no official prerequisite for this course. The key requirements are:

1. **Willingness to learn** materials that may be outside your current field/discipline. This course includes content from Computer Science, Statistics, and other disciplines.
2. **Passion for learning new technologies** without apprehension, particularly given the rapid evolution of AI tools.
3. **Commitment to apply** what is learned in your own research.
4. **True to yourself** since all things can be exposed in AI age.

4 Learning Objectives

Upon completion of this course, students will be able to:

1. **Understand** the foundational principles of knowledge mining and its role in contemporary research, including the emerging AI for Science (AI4S) paradigm.
2. **Apply** advanced techniques in data mining, machine learning, and AI, including reasoning models, RAG architectures, and agentic AI workflows.

3. **Develop** practical skills in AI-assisted tools for data analysis, text mining, and visualization, with emphasis on maintaining scientific rigor.
4. **Design and implement** RAG pipelines and knowledge graph systems for research applications.
5. **Evaluate** AI-generated outputs critically, understanding limitations and potential for hallucination.
6. **Investigate** ethical implications of knowledge mining and AI applications in research.
7. **Collaborate effectively** with AI systems as research partners across scaffolded interaction levels.

5 Required Texts and Readings

5.1 Primary Texts

Text	Notes
Bowen, J. A., & Watson, C. E. (2024). <i>Teaching with AI: A Practical Guide to a New Era of Human Learning</i> . Johns Hopkins University Press.	Available electronically at McDermott Library
Breiman, L. (2001). Statistical modeling: The two cultures. <i>Statistical Science</i> , 16(3), 199-231.	Available electronically
Daoud, A. & Dubhashi, D. (2020). Statistical modeling: The three cultures. arXiv:2012.04570	Link

5.2 New Required Readings (2026)

Text	Topic	Notes
Poldrack, R. A. et al. (2025). Ten Simple Rules for AI-Assisted Coding in Science. arXiv:2510.22254	AI-Assisted Coding	Link

Text	Topic	Notes
Kamalov, F. et al. (2025). Evolution of AI in Education: Agentic Workflows. arXiv:2504.20082	Agentic AI	Link
Nature (2025). AI for Science 2025 Collection.	AI4S Paradigm	Available at UTD Library

5.3 Recommended Reading

5.3.1 Data Mining

- Han, J., Kamber, M., & Pei, J. (2012). *Data Mining: Concepts and Techniques* (3rd ed.). O'Reilly. (HKP)

5.3.2 Statistical Learning and Machine Learning

- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *Introduction to Statistical Learning with Applications in R* (2nd ed.). Springer. (ISLR) [Website](#)
- Kuhn, M. & Silge, J. (2021). *Tidy Modeling with R*. [Online](#) (KS)

5.3.3 Text Mining and NLP

- Silge, J. & Robinson, D. (2022). *Text Mining with R: A Tidy Approach*. [Online](#) (SR)

5.3.4 AI and Data Science

- Ho, K. (2024). *Data Programming with GenAI*. [Online](#)
- Molnar, C. (2019). *Interpretable Machine Learning*. [Online](#)

5.3.5 For Social Scientists

- Grimmer, J., Roberts, M. E., & Stewart, B. M. (2021). Machine Learning for Social Science: An Agnostic Approach. *Annual Review of Political Science*, 24(1), 395–419.
- Attewell, P. A. & Monaghan, D. B. (2015). *Data Mining for the Social Sciences: An Introduction*. UC Press.

5.3.6 For Economists

- Athey, S. (2018). The impact of machine learning on economics. In *The Economics of Artificial Intelligence* (pp. 507-547). University of Chicago Press.

- Varian, H. R. (2014). Big data: New tricks for econometrics. *Journal of Economic Perspectives*, 28(2), 3-28.

6 Course Structure

6.1 Thematic Organization

The course is organized around four major themes:

Phase	Weeks	Focus Area
I. Foundations	1–4	Data science principles, AI paradigms, two/three cultures
II. Methods	5–8	ML techniques, NLP, LLMs, RAG architectures
III. AI Integration	9–13	AI-assisted coding, synthetic data, evaluation metrics
IV. Applications & Ethics	14–16	Ethics, future directions, project presentations

6.2 AI Interaction Scaffolding

Students will progress through three levels of AI collaboration:

Phase	Weeks	AI Role	Student Responsibility
Foundation	1–5	AI as Assistant	Human initiates; AI responds to specific queries
Collaboration	6–10	AI as Collaborator	Iterative dialogue; shared agency in problem-solving
Delegation	11–16	AI as Delegatee	Human sets goals; AI executes with oversight

7 Grading and Requirements

7.1 Grade Structure

Component	Weight	Description
Participation & AI Collaboration Log	10%	Active participation + documented AI interactions

Component	Weight	Description
Assignments (posted on website)	10%	Applied exercises demonstrating techniques
Project Proposal	20%	Original research design with AI integration plan
Progress Reports (2)	30%	Iterative development with documented AI collaboration
Final Project & Presentation	30%	Original research with methodology documentation

7.2 Personal Website Requirement

Each student must create and maintain a website hosted on **GitHub**. All assignments and class materials should be posted on your website before the next class.

 Recommended Method

Use **Quarto** for building your GitHub website.

7.3 Project Requirements

The final project must be **20–25 pages** in length. Requirements include:

- **Originality:** Topic, data, or method(s) must be originally designed and unpublished
- **Consultation:** Proposal must be discussed with instructor and approved
- **Replication:** Acceptable with strong justification (e.g., dissertation research)
- **AI Documentation:** Full documentation of AI tool usage required

 Important

No Kaggle project/data replications. Period. Projects must demonstrate original contribution. Python can be used but only to supplement R-based analysis. If you insist to use Python, do extra steps to convert to R (e.g., reticulate, rpy2, etc.) and document the process. Notify the instructor.

7.4 Assignment Guidelines

Exercises apply data programming techniques in:

- Data modeling and visualization
- Exploratory data analysis

- AI-assisted code development and validation

Students are encouraged to develop techniques beyond class materials while maintaining scientific rigor.

8 Tentative Schedule

i Note

This schedule is subject to change. Updates will be announced in class and on Teams.

8.1 Part I: Foundations (Weeks 1–4)

Week	Date	Topic	Due
1	Jan 20	Introduction to Knowledge Mining & AI4S Paradigm	
2	Jan 27	Data Science Foundations: Two/Three Cultures	
3	Feb 3	Latest Trends: RAG, Agentic AI, Knowledge Graphs	
4	Feb 10	Machine Learning for Knowledge Mining: Reasoning Models	

8.2 Part II: Methods & Techniques (Weeks 5–8)

Week	Date	Topic	Due
5	Feb 17	Causal and Predictive Modeling	
6	Feb 24	Text Mining, NLP, and Multimodal Knowledge Extraction	
7	Mar 3	LLMs, Reasoning Models, and Agentic AI Systems	Proposal
8	Mar 10	AI for Research and Knowledge Discovery: RAG Pipelines	

8.3 Part III: AI Integration (Weeks 9–13)

Week	Date	Topic	Due
9	Mar 17	<i>Spring Break - No Class</i>	
10	Mar 24	AI-Assisted Data Collection and Coding	
11	Mar 31	Synthetic Data and Silicon Samples	Progress 1
12	Apr 7	Evaluation Metrics, Hallucination Detection, Model Selection	
13	Apr 14	Advanced Applications: Multi-Agent Systems for Research	

8.4 Part IV: Ethics & Applications (Weeks 14–16)

Week	Date	Topic	Due
14	Apr 21	Data Ethics in Knowledge Mining	
15	Apr 28	Future Directions in Knowledge Mining	Progress 2
16	May 5	Final Project Presentations	Final

9 Course Policies

9.1 Participation

Full attendance is **required and imperative**. Participation includes:

- Full preparation for class including completion of assigned readings
- Active involvement in class discussion
- Documented AI collaboration (AI Collaboration Log)

Only medical emergencies (with documentation) constitute excused absences. Contact instructor in first week for special medical accommodations.

9.2 Generative AI Policy

i AI Usage Policy

Generative AI is permitted in specific contexts and **requires acknowledgment**.

When using AI tools on assignments, include an appendix with:

- a. **Full exchange:** The entire conversation, highlighting relevant sections
- b. **Tool specification:** Which AI tools were used (e.g., “Claude 3.5 Sonnet via API”)
- c. **Usage explanation:** How AI was used (idea generation, code drafting, text refinement, etc.)
- d. **Rationale:** Why AI tools were used (efficiency, overcoming blocks, experimentation, etc.)

9.2.1 AI Collaboration Log

In addition to assignment appendices, maintain a running **AI Collaboration Log** documenting:

- Patterns of effective AI interaction
- Failed approaches and lessons learned
- Evolution of your AI collaboration skills

⚠ Warning

AI tools should be used **wisely and reflectively** to deepen understanding, not to bypass learning. The goal is developing adaptive collaboration skills for an AI-augmented research future.

9.3 Document Guidelines

All documents must adhere to:

- **Format:** PDF only (no Word or Google Docs)
- **Naming:** `lastname_documenttype.pdf` (e.g., `smith_proposal.pdf`)
- **Paper:** US Letter size
- **Margins:** 1 inch on all sides
- **Font:** 12-point
- **Style:** APSA format ([Style Manual](#))
- **Submission:** Email attachment (no cloud drive links)

9.4 Class Q&A

Piazza is used for class discussion. Before emailing, check/post on Piazza.

- Enroll: <https://piazza.com/utdallas/spring2026/epps6323>

- Access code: **1502** (classroom number)

9.5 Software and Tools

9.5.1 Required Tools

- R and RStudio
- Python 3.x
- Git and GitHub account
- Quarto

9.5.2 AI Tools

Students should have access to at least one of the following:

- Claude (Anthropic)
- ChatGPT (OpenAI)
- GitHub Copilot (recommended for coding)

9.5.3 Style Guide

Follow the [Tidyverse Style Guide](#) for R code.

10 University Policies

10.1 Comet Creed

“As a Comet, I pledge honesty, integrity, and service in all that I do.”

This creed was voted on by the UT Dallas student body in 2014.

10.2 Academic Support Resources

See <http://go.utdallas.edu/academic-support-resources> for university academic support resources.

10.3 UT Dallas Syllabus Policies and Procedures

The information at <http://go.utdallas.edu/syllabus-policies> constitutes the University’s policies and procedures segment of this syllabus.

Please review catalog sections regarding credit/no credit grading and withdrawal policies.

i Disclaimer

The descriptions and timelines in this syllabus are subject to change at the discretion of the instructor. Changes will be communicated via class announcements and Piazza.

11 Appendix: Key Concepts in 2026 Knowledge Mining

11.1 Agentic AI

AI systems that can reason, plan, use tools, and execute multi-step tasks autonomously. Key patterns include:

- **Tool Use:** AI accessing external tools and APIs
- **Reflection:** AI evaluating and improving its own outputs
- **ReAct:** Reasoning and Acting in interleaved steps
- **Planning:** Breaking complex tasks into sub-goals
- **Multi-Agent Collaboration:** Multiple AI agents working together

11.2 Retrieval-Augmented Generation (RAG)

Architecture that enhances LLM outputs by retrieving relevant information from external knowledge bases. Components:

- **Indexing:** Processing and storing documents as embeddings
- **Retrieval:** Finding relevant documents for a query
- **Generation:** Producing outputs grounded in retrieved context

11.3 Knowledge Graphs

Structured representations of knowledge using nodes (entities) and edges (relationships). In 2026, increasingly integrated with:

- Multimodal data (text, images, video)
- Vector databases for semantic search
- LLMs for knowledge extraction and reasoning

11.4 AI for Science (AI4S)

The convergence of AI innovation with scientific research, representing a transformative research paradigm that integrates:

- Experimental science (empirical induction)
- Theoretical science (theoretical modeling)

- Computational science (simulation)
- Data-intensive science (pattern recognition)
- AI-driven discovery (autonomous hypothesis generation)